

Problem statement

A real estate agent want help to predict the house price for regions in USA.He gave us the dataset to work on to use Linear Regression model.Create a model that hepls to determine

Data collection

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing dataset

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\H.csv")
data
```

Out[2]:

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price | Address |
|------|---------------------|------------------------------|---------------------------------------|---------------------------------------|--------------------|--------------|---------------------------------------------------------|
| 0 | 79545.45857 | 5.682861 | 7.009188 | 4.09 | 23086.80050 | 1.059034e+06 | 208 Michael Ferry Apt. 674\nLaurabury, NE 3701... |
| 1 | 79248.64245 | 6.002900 | 6.730821 | 3.09 | 40173.07217 | 1.505891e+06 | 188 Johnson Views Suite 079\nLake Kathleen, CA... |
| 2 | 61287.06718 | 5.865890 | 8.512727 | 5.13 | 36882.15940 | 1.058988e+06 | 9127 Elizabeth Stravenue\nDanieltown, WI 06482... |
| 3 | 63345.24005 | 7.188236 | 5.586729 | 3.26 | 34310.24283 | 1.260617e+06 | USS Barnett\nFPO AP 44820 |
| 4 | 59982.19723 | 5.040555 | 7.839388 | 4.23 | 26354.10947 | 6.309435e+05 | USNS Raymond\nFPO AE 09386 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 4995 | 60567.94414 | 7.830362 | 6.137356 | 3.46 | 22837.36103 | 1.060194e+06 | USNS Williams\nFPO AP 30153-7653 |
| 4996 | 78491.27543 | 6.999135 | 6.576763 | 4.02 | 25616.11549 | 1.482618e+06 | PSC 9258, Box 8489\nAPO AA 42991- 3352 |

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price | Address |
|------|---------------------|------------------------------|---------------------------------------|---------------------------------------|--------------------|--------------|---------------------------------------------------|
| 4997 | 63390.68689 | 7.250591 | 4.805081 | 2.13 | 33266.14549 | 1.030730e+06 | 4215 Tracy Garden Suite 076\nJoshualand, VA 01... |
| 4998 | 68001.33124 | 5.534388 | 7.130144 | 5.44 | 42625.62016 | 1.198657e+06 | USS Wallace\nFPO AE 73316 |
| 4999 | 65510.58180 | 5.992305 | 6.792336 | 4.07 | 46501.28380 | 1.298950e+06 | 37778 George Ridges Apt. 509\nEast Holly, NV 2... |

5000 rows × 7 columns

head

In [3]:

```
# to display first 8 dataset values
data.head(8)
```

Out[3]:

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price | Address |
|---|---------------------|------------------------------|---------------------------------------|---------------------------------------|--------------------|--------------|---------------------------------------------------|
| 0 | 79545.45857 | 5.682861 | 7.009188 | 4.09 | 23086.80050 | 1.059034e+06 | 208 Michael Ferry Apt. 674\nLaurabury, NE 3701... |
| 1 | 79248.64245 | 6.002900 | 6.730821 | 3.09 | 40173.07217 | 1.505891e+06 | 188 Johnson Views Suite 079\nLake Kathleen, CA... |
| 2 | 61287.06718 | 5.865890 | 8.512727 | 5.13 | 36882.15940 | 1.058988e+06 | 9127 Elizabeth Stravenue\nDanieltown, WI 06482... |
| 3 | 63345.24005 | 7.188236 | 5.586729 | 3.26 | 34310.24283 | 1.260617e+06 | USS Barnett\nFPO AP 44820 |
| 4 | 59982.19723 | 5.040555 | 7.839388 | 4.23 | 26354.10947 | 6.309435e+05 | USNS Raymond\nFPO AE 09386 |
| 5 | 80175.75416 | 4.988408 | 6.104512 | 4.04 | 26748.42842 | 1.068138e+06 | 06039 Jennifer Islands Apt. 443\nTracyport, KS... |
| 6 | 64698.46343 | 6.025336 | 8.147760 | 3.41 | 60828.24909 | 1.502056e+06 | 4759 Daniel Shoals Suite 442\nNguyenburgh, CO ... |
| 7 | 78394.33928 | 6.989780 | 6.620478 | 2.42 | 36516.35897 | 1.573937e+06 | 972 Joyce Viaduct\nLake William, TN 17778-6483 |

info

In [4]:

```
# to identify missing values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

describe

In [5]:

```
# to display summary of the dataset
data.describe()
```

Out[5]:

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price |
|--------------|---------------------|------------------------|---------------------------------|------------------------------------|--------------------|--------------|
| count | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 | 5.000000e+03 |
| mean | 68583.108984 | 5.977222 | 6.987792 | 3.981330 | 36163.516039 | 1.232073e+06 |
| std | 10657.991214 | 0.991456 | 1.005833 | 1.234137 | 9925.650114 | 3.531176e+05 |
| min | 17796.631190 | 2.644304 | 3.236194 | 2.000000 | 172.610686 | 1.593866e+04 |
| 25% | 61480.562390 | 5.322283 | 6.299250 | 3.140000 | 29403.928700 | 9.975771e+05 |
| 50% | 68804.286405 | 5.970429 | 7.002902 | 4.050000 | 36199.406690 | 1.232669e+06 |
| 75% | 75783.338665 | 6.650808 | 7.665871 | 4.490000 | 42861.290770 | 1.471210e+06 |
| max | 107701.748400 | 9.519088 | 10.759588 | 6.500000 | 69621.713380 | 2.469066e+06 |

columns

In [6]:

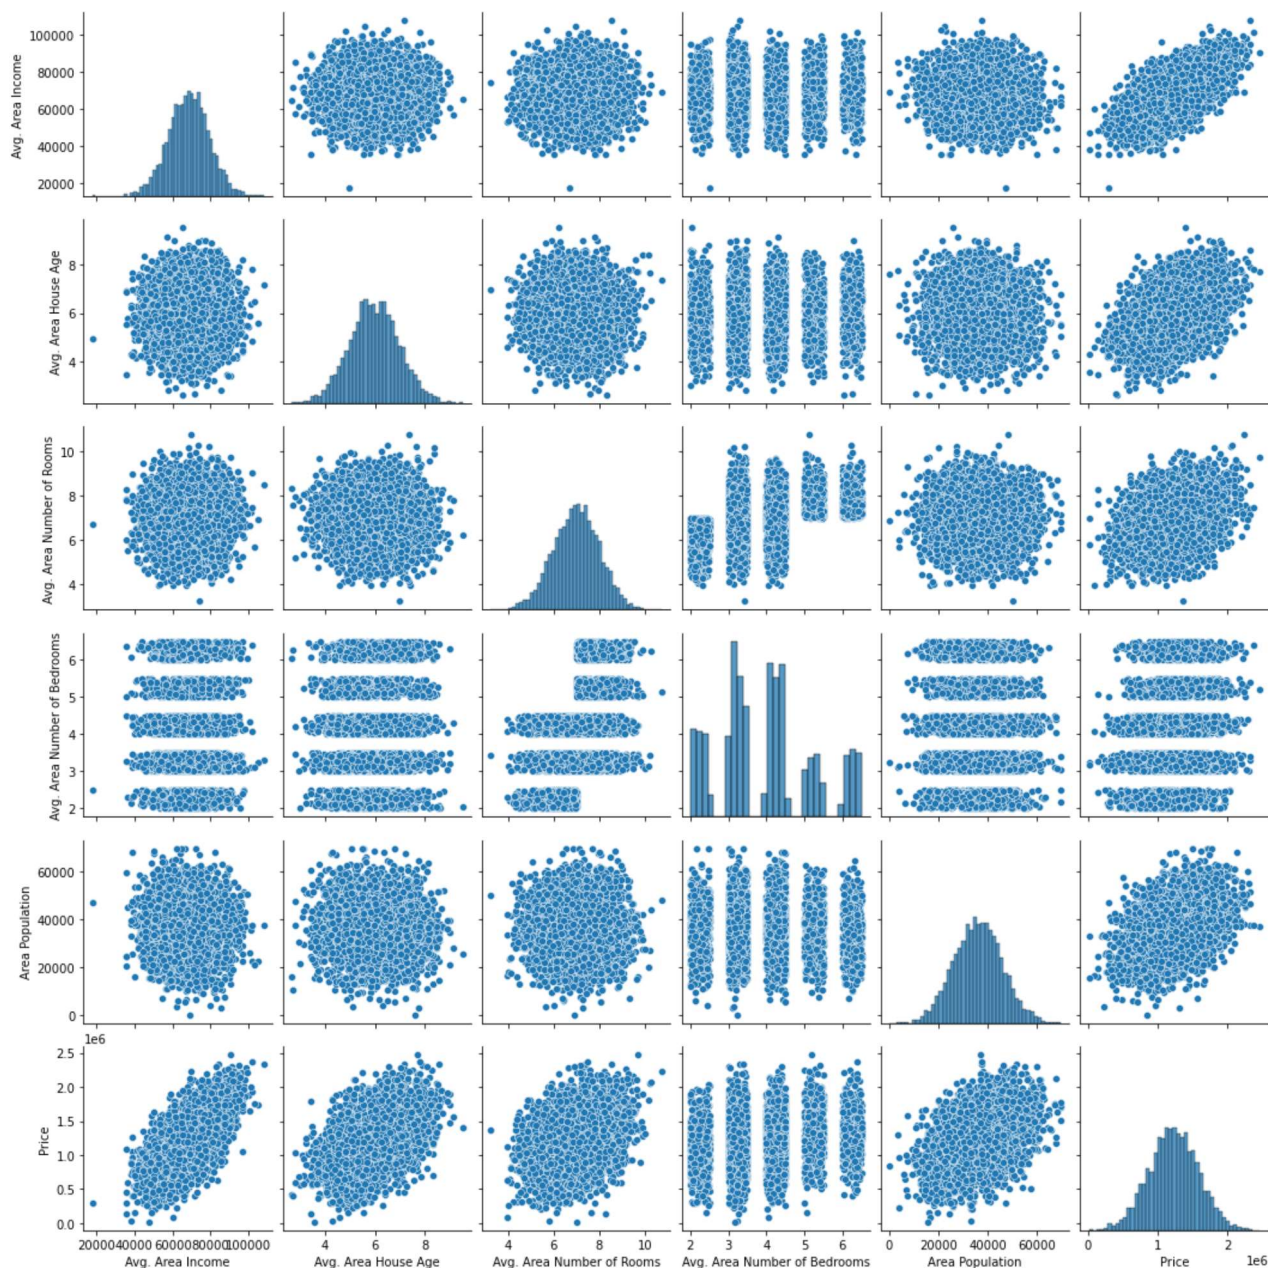
```
# to display headings of the dataset
data.columns
```

Out[6]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
dtype='object')

EDA and Visualization

```
In [7]: sns.pairplot(data)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x248b465f430>
```



distribution plot

```
In [8]: sns.displot(data[["Area Population"]])
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x248b465f910>
```



correlation

```
In [9]: da=data[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
               'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address']]  
sns.heatmap(da.corr())
```

Out[9]: <AxesSubplot:>



To train the model-Model Building

we are going to train Linear Regression model;we need to split out data into two variables x and y where x is independent variable and y is dependent variable on x (output) we could ignore address column as it is not required for our model

```
In [10]: x=da[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
            'Avg. Area Number of Bedrooms', 'Area Population']]
        y=da['Price']
```

```
In [11]: # to split my dataset into training and test data
        from sklearn.model_selection import train_test_split
        x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [12]: from sklearn.linear_model import LinearRegression
        lr= LinearRegression()
        lr.fit(x_train,y_train)
```

```
Out[12]: LinearRegression()
```

```
In [13]: print(lr.intercept_)
```

```
-2637937.5603702217
```


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
        coeff
```

```
Out[14]:
```

| | Co-efficient |
|-------------------------------------|---------------|
| Avg. Area Income | 21.610966 |
| Avg. Area House Age | 165415.423387 |
| Avg. Area Number of Rooms | 120901.707576 |
| Avg. Area Number of Bedrooms | 920.427999 |
| Area Population | 15.232640 |

```
In [15]: prediction=lr.predict(x_test)
        plt.scatter(y_test,prediction)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x248b7791970>
```

```
In [16]: print(lr.score(x_test,y_test))
```

```
0.9157772232713607
```

```
In [17]: lr.score(x_train,y_train)
```

```
Out[17]: 0.9188884356053647
```

Ridge regression

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

```
Out[21]: 0.9157728251515059
```

```
In [22]: rr.score(x_train,y_train)
```

```
Out[22]: 0.9188855005670066
```

LaSSO regression

```
In [26]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
la.score(x_train,y_train)
```

```
Out[26]: 0.9188884339089254
```

```
In [29]: la.score(x_test,y_test)
```

Out[29]: 0.9157768184155307

In []: